

SlickChair

Olivier Blanvillain
École Polytechnique Fédérale de Lausanne
1015 Lausanne, Switzerland
olivier.blanvillain@epfl.ch

ABSTRACT

Ut in dolor et magna tincidunt mattis. Proin id pulvinar arcu. Donec ac turpis consectetur, dignissim eros at, mollis orci. Nunc sed tincidunt justo, eu dapibus risus. Vestibulum nisi mi, tempus nec cursus at, accumsan vitae metus. Cras mattis, velit ut convallis lacinia, metus enim rutrum sapien, quis egestas ante ipsum ut lacus. Aliquam erat volutpat. Mauris vitae commodo nisi. Nunc iaculis, enim vulputate cursus interdum, sapien libero sodales diam, viverra molestie diam tellus eu felis. Quisque gravida porttitor vulputate. Duis nec neque facilisis, porttitor ante id, molestie sem. Vivamus pellentesque venenatis est, ut consequat arcu tempus a. Phasellus ullamcorper nunc vel rhoncus suscipit. Curabitur commodo ornare accumsan. Mauris vitae lorem arcu. Mauris vel turpis mi. Fusce faucibus congue ante, eu gravida sem. Cum sociis natoque penatibus et magnis dis.

1. INTRODUCTION

Peer reviewing is a central process in the organisation of scientific conferences. It involves multiple interactions between it's participants: *authors* share their work with the *program committee* which is then in charge of reviewing these submissions. A *program chair* is designated to coordinate the process and decide on final submission acceptance. While this peer review process could be implemented by simply exchanging emails between participants, manually gathering submissions, assigning them to program committee members, aggregating the reviews and finally notifying the authors requires substantial efforts. The use of a dedicated software, called conference management system, can greatly simplify this process.

Over the last few years, numerous web based conference management system have been developed. A recent comparative study [14] shows that EasyChair [5] is by large the most popular platform, having been used in about 68% of conferences organized with online conference management system. The following four systems in term of number of organized conferences are EDAS [6] with 8.5%, Open Conference Systems [12] with 6%, START V2 [17] with 5.7% and ConfTool [2] with 5.3%. However, out of these five systems, only Open Conference Systems is open-source. EasyChair and ConfTool offer free licenses of their restricted versions and START V2 and EDAS are only available as a commercial product.

The importance and confidentiality of the data manipulated by a conference management system imply that security a major concern when working with such system. Closed

source solutions are often only available as hosted services, therefore requiring conference organisers to trust the company behind the product not only for the quality of the code, but also for the robustness of the infrastructure and the respect of data privacy. The open-source solutions we considered where not providing satisfactory levels of security. For example, Open Conference Systems sends a copy of the user passwords by email as plain text once the registration is completed. HotCRP [10], another open-source conference management system, has a similar flaw: it sends login links to users with the password as part of the url.

We present SlickChair, an open-source conference management system written in Scala. Build with the Play framework and the Slick database access library, SlickChair provides a highly flexible and extensible solution to manage peer review processes. Our contributions are in particular:

- The plan

2. OVERVIEW OF SLICKCHAIR

In this section gives an overview of functionalities of SlickChair. We first discuss how users.. (login, actors, phases (work-flow))

2.1 User authentication

The first contact between a system and its users is often via an authentication interface. Most conference management systems require users to create new accounts, which usually implies filling a form, receiving a validation email and memorizing a new password. Recent web browser implement mechanisms to reduced the pain of filling form and memorizing passwords, but these solutions are usually limited to the use of a single computer.

In SlickChair, we address this common problem by supporting three types of authentication mechanisms:

- Authentication via Facebook account
- Authentication via Google account
- Authentication via email address

The implementation of Facebook and Google login uses the OAuth 2.0 protocol [9], and is provided by the Secure-Social authentication module for Play Framework [16].

In addition to associate each visitor with a unique identity, an objective of authentication is to get a valid email address to contact the user. This way, we avoid any chances of having typo in the email. When users login with Facebook or Google account, we trust the email address obtained via the OAuth protocol, given that both networks required their users to confirm their email address when they created their account. In the case of login via email, the process is obviously more complex. After providing his email address, a new user has to open his email client follow a validation link. From here, the user is asked his first name, last name, and a new password (of minimum 8 characters) before completing the registration. Passwords are hashed using the bcrypt algorithm [15] and stored in the database. SlickChair also gives its users the usual options to change their password or recover their account in case of forgotten password. Although authentication via email address add complexity to the system, it is appreciated by users that have separated email accounts for their personal and professional life.

In SlickChair, we identify each user by a single email address. Some other conference management systems provide the ability to link multiple email addresses to a single identity and to multiple merge accounts into one. We believe that such functionality can sometimes be confusing, and adds a lot of complexity to the system. As a side note, this design makes it possible for a user to close the Facebook or Google account he used to login, and still claim his identity by going through the process of logging in using the corresponding email address.

2.2 SlickChair interfaces

Building SlickChair, our focused was on creating a flexible and extensible system rather than offering customization options via configuration. As a system becomes more complete, it's complexity is likely to go up, and maintaining and extending the system becomes harder. SlickChair provides the essential¹ components to run an online peer review process, listed below as the user interfaces of available for each user role:

- Author
 - Make new submissions
 - Edit submissions
 - See reviews
- Program Committee Member
 - Bid on submissions
 - Write reviews
 - Write comments
- Program Chair
 - Assign submissions for review
 - Decide on submissions acceptance

¹The authors of [4] identified nine *typical* functionalities offered by conference management systems to run online peer review processes, which roughly correspond to the functionalities provided by SlickChair interfaces.

- Change user roles
- Send emails and change conference phases

The *Change user roles* interface allows a program chair to design certain users as being co-chairs or program committee members. While it is also possible to define all program committee members and chairs in a configuration file before setting up a SlickChair instance, using this *change user roles* interface might be preferred to allow each user to chose his favourite login method. This accounts for the typical situation of users that forward messages send to their professional email addresses to other services such as Gmail.

The *Send emails and change conference phases* interface is related to the conference workflow, discussed in the following subsection. *Assign submissions for review* is discussed in detail ???. Other interfaces are quite simple and wont be described for brevity.

2.3 Workflow

In order to coordinate the overall conference peer review process, we organised the course of this process as a chronological sequence on phases. Each phase corresponds to the set of interfaces enabled during the during the particular time frame of the phase. We identified the following seven phases that may correspond to the workflow of a small conference: *Setup*, *Submission*, *Bidding*, *Assignment*, *Review*, *Notification*, *Finished*.

In addition to the configuration of enabled interfaces, each phase is defined with a function to generate emails, and function to generate an optional warning:

```
case class Phase(
  configuration: Configuration,
  emails: Database => List[Email],
  warning: Database => Option[String]
)
```

SlickChair uses the `emails` function before transitioning to a given phase to help the program chair in the redaction of notification emails, by suggesting appropriate recipients and giving a template for the body of the message. If a warning is returned by the second function, it will be display to the program chair before he confirms the change of phase and the sending of notifications.

Conferences of different sizes and topics might want to use different workflows. In the current stage of the project, such configuration has be done has source code level. As an example, we will show the changes needed to add a new component to the conference workflow. Suppose that some program committee members are careless about their reviewing responsibilities and do not sent respect the delays set by the program chair. The program chair might to send a reminder to the program committee members that have not yet completed their reviews. This could be implemented by adding the following phase between *Review* and *Notification*:

```
Phase(
  Configuration("Review Reminder",
    pcmemberCanReview=true,
```

```

    pcMemberCanComment=true,
    chairCanDecideOnAcceptance=true),
  { db => Email(
    lateReviewerEmails(db),
    "Reminder: review deadline",
    "Dear Program Committee Member, ...")},
    noWarning
  )

```

Where `noWarning` a dummy function that never returns any warning, and `lateReviewerEmails` is a function returning the email addresses of all late reviewers. In order to compute the late reviewers, we need to use three tables of the database: the `assignments` and `reviews` tables are relations between program committee members and submissions, and the `persons` table contains personal information of SlickChair users.

```

def lateReviewerEmails(db: Database) = {
  val assignmentPairs = db.assignments
    .filter(_._value)
    .map(a => (a.personId, a.paperId))
  val reviewsPairs = db.reviews
    .map(r => (r.personId, r.paperId))
  (assignmentPairs diff reviewsPairs)
    .join(db.persons).on(_._1 is _._id)
    .map(_._2.email).list(db.s)
}

```

This function illustrates the use of the Slick database query library. Except for `.join().on()` and `.list(db.s)`, the code would be identical if manipulating Sets from the Scala collections. In reality, the entire body of this function will be compiled into a SQL query. The `.join().on()` allows to join two tables on certain columns, and returns pairs of matching rows. `.list()` wraps up the query definition and returns the result of execution as a list of Scala objects.

3. DATA MODEL

One of SlickChair's feature is its ability to log all the actions and events that occurred in the system. To do so, a traditional approach could use the logging functionalities build into the Play framework. Concretely, this would corresponds to appending a timestamp and a messages to a text file whenever an interesting event occurs. We took an alternative approach which consists in memorizing every change at the level of the database.

The idea of append only databases is not new, but it is becoming more and more relevant as the prices of storage go down. The Datomic database [3] which recently entered the market is build around this idea of never altering or deleting data. Instead, changes are made by adding new version of the existing data or marking it as being invalid. However, Datomic is a proprietary software and does not fit well with the open-source platform SlickChair is build upon.

In order to suite our needs, we build a small layer on top of Slick to manipulate versioned data. The semantic of the API presented in this section are inspired by Datomic's Clojure API [1].

3.1 Database as a value

One of the core concept of functional programming the manipulation of immutable data. But how could a database be seen as immutable? Instead if manipulating directly the live, mutating database, we define a `Database` to be a view of the data on at a particular time date:

```

case class Database(
  val date: DateTime, val s: Session) {
  val persons = table[Person]
  val personRoles = table[PersonRole]
  val papers = table[Paper]
  val paperAuthors = table[PaperAuthor]
  val paperDecisions = table[PaperDecision]
  val comments = table[Comment]
  val reviews = table[Review]
  val files = table[File]
  val emails = table[Email]
  val bids = table[Bid]
  val assignments = table[Assignment]
  val configurations = table[Configuration]
}

```

Here, manipulations done on this object, such as the function `lateReviewerEmails` define in subsection 2.3, will see the world as it was at time `date`. Hence, calling this function with the same argument will always return the same result: the query is a pure function and the `Database` is a value.

A `Connection` allows to retrieve the value of the current `Database`, and insert new elements in the database. The `insert` method returns two `Database` values, value just before and just after the insertion.

```

case class Connection(s: Session) {
  def currentDatabase(): Database
  def insert(x: List[Model[_]]):
    (Database, Database)
}

```

The last definition of this API is the `Model` trait, which enforces that the elements stored in tables are case classes defined with a `metadata` field with the appropriate information:

```

case class Id[M](value: UUID)

type Meta[M] = (Id[M], DateTime, String)

trait Model[M] {
  this: M { def metadata: Meta[M] } =>
  val (id, updatedAt, updatedBy) = metadata
}

```

This API takes full advantage of Scala and Slick type-checking capabilities. Thanks to the type parameter on `Id`, the compiler will detect errors such as miss using the `Id` of a person as the `Id` of a submission.

3.2 Implementation

Under the hood, SlickChair create multiple database records for a given *natural key*. When using the `Database` object, the user defined query is composed with a temporal filter that extracts the most recent records that where created

before the `date` timestamp of the `Database`. This corresponds to the *Type 2 slowly changing dimension management methodology*, as described in [11].

This approach has the advantage to not require any update, new records are simply appended to tables and only become visible for the newly accessed `Database` values. It also minimises the number of tables, which is beneficial as tables directly manipulated by the user via Slick's domain-specific language. We developed this API with H2 and PostgreSQL, but as the implementation uses Slick's generic `JdbcDriver`, it should be compatible with all Slick's supported database systems.

Finally, we should note that in its current stage, our *database as a value* abstraction does not allow the expression of transactions. Our timestamp based implementation is by nature atomic: the temporal filter used on every table ensures that the database will be viewed either before or after an insertion. However it is currently impossible to express a *transaction*, which would mean enforce that the database is not modified between the time when its value is queried and the insertion are effective. Fortunately SlickChair does not require such semantic.

4. SUBMISSION-REVIEWER ASSIGNMENT

One of the main responsibility of the program chair is the assignment of submissions to program committee members. While this could reasonably be done manually for a small number of submissions, the task quickly becomes complex as the number of submissions goes up. In fact, several mathematical formulations are known to be NP-Hard. We will see how we formulated the problem using the OscaR operational research library, which implements search patterns to explore the space of all possible assignments under limited time constraints.

To be fair to all authors, submissions usually receive the same number of reviews, and this work has to be well distributed among program committee members so that no one is overloaded. Moreover, program committee members might have conflicts of interests and varying levels of knowledge depending on the topics. Given these constraints, one can associate a numeric value to each submission-reviewer assignment in order to create a total order on all the valid assignments.

4.1 Mathematical formulation

This problem can be seen as a generalization of the stable marriage problem. On one side, the reviewers emitted preferences for the submissions in the form of bids, called the *interest preference*. On the other side, submissions can be seen as being willing to be matched with expert reviewers, the *expertise preference*. In the case of a peer review system, the problem is generalized in the following ways:

- *Polyamory*: submissions and reviews are matched many-to-many instead of one-to-one.
- *Indifference*: interest and expertise preference contain ties between instead of being total ordered.

- *Incomplete lists*: conflict of interest forbids certain assignment, removing element from the candidate lists.

While each individual variant can be solved in polynomial time, the combination of the three generalizations is NP-complete [8]. An approximation algorithm to maximize fairness among the review is discussed in [7], but it revolves around repetitively solving linear programmes, which is not a practical solution.

4.2 OscaR formulation

OscaR is a Scala library for solving Operations Research problems [13]. Among other techniques, OscaR allows to formulate problems using constraint programming. Given a matrix `bids`, a list of `conflicts`, and four integers `nPapers`, `nReviewers`, `nReviewPerPaper` and `nReviewsPerReviewer`, the following program searches for an assignment maximizing the sum of reviewer satisfaction:

```
val m = makeMatrix(nReviewers, nPapers, 0 to 1)

0 until nPapers foreach { j =>
  add(sum(m col j) == nReviewPerPaper) }
0 until nReviewers foreach { i =>
  add(sum(m row i) == nReviewsPerReviewer) }
conflicts foreach {
  add(m(_, _) == 0) }

maximize(weightedSum(bids, m)) search {
  binaryMaxDegree(m.flatten.toSeq) }
onSolution { return m } start()
```

In this code, the `add` function adds a constraint to the definition problem. In this order, it is used to enforce that all submissions have the same number of review, all reviewers have the same of review, and no conflicting assignment is made. The `maximize()` `search {}` and `onSolution {} start()` are OscaR constructs to define the objective function, the search heuristic and start the resolution. Some preparation is required to assemble the data in an appropriate form, which notably includes the addition of dummy submissions such that repartition of reviews is perfectly balanced among reviewers. Dummy papers are assigned with the maximum bid, which according to the experimental evaluation of [7] leads to better results.

5. CONCLUSION AND FUTURE WORK:

- Macros (remove boilerplate on the database API)
- Scala.js

6. REFERENCES

- [1] ClojureDatomic Clojure API. Available at: <http://docs.datomic.com/clojure/>.
- [2] ConfTool Website. Available at: <http://www.conftool.net>.
- [3] Datomic. Available at: <http://www.datomic.com/>.
- [4] N. Di Mauro, T. M. A. Basile, and S. Ferilli. GRAPE: An expert review assignment component for scientific conference management systems. In *Proceedings of the 18th International Conference on Innovations in Applied Artificial Intelligence*, 2005.

- [5] Easy Chair Website. Available at: <http://edas.info/doc>.
- [6] EDAS Website. Available at: <http://www.easychair.org>.
- [7] N. Garg, T. Kavitha, A. Kumar, K. Mehlhorn, and J. Mestre. Assigning papers to referees, 2010.
- [8] J. Goldsmith and R. H. Sloan. The AI conference paper assignment problem. In *Proc. AAAI Workshop on Preference Handling for Artificial Intelligence, Vancouver*, 2007.
- [9] D. Hardt. The OAuth 2.0 authorization framework. RFC 6749, RFC Editor, 2012.
- [10] HotCRP Website. Available at: <http://read.seas.harvard.edu/~kohler/hotcrp/>.
- [11] R. Kimball and M. Ross. *The Data Warehouse Toolkit: The Complete Guide to Dimensional Modeling*. John Wiley & Sons, Inc., New York, NY, USA, 2nd edition, 2002.
- [12] Open Conference System Website. Available at: <http://pkp.sfu.ca/?q=ocs>.
- [13] Oscala Team. Oscala: Scala in OR, 2014. Available at: <http://bitbucket.org/oscarlib/oscar>.
- [14] L. Parra, S. Sendra, S. Ficarelli, and J. Lloret. Comparison of online platforms for the review process of conference papers. In *The Fifth International Conference on Creative Content Technologies*, 2013.
- [15] N. Provos and D. Mazières. A future-adaptive password scheme. In *Proceedings of the Annual Conference on USENIX Annual Technical Conference, ATEC '99*, 1999.
- [16] Secure Social. Available at: <http://securesocial.ws/>.
- [17] START V2 Website. Available at: <http://www.softconf.com/>.